



ErSE222: Machine learning in Geoscience

March. 2nd, 2025

Tariq Alkhalifah and Omar Saad





SGD with momentum


Idea of momentum (from Polyak and Nesterov in 60'): use update that is an exponentially decaying moving average of the past gradients.


Velocity Momentum term [0,1)

↓ ↙

$$\mathbf{v}_{i+1} = \gamma \mathbf{v}_i - \mathbf{g}_{i+1} = \gamma \mathbf{v}_i - \frac{\alpha}{N_b} \sum_{j=1}^{N_b} \nabla \mathcal{L}_j$$

$$\boldsymbol{\theta}_{i+1} = \boldsymbol{\theta}_i + \mathbf{v}_{i+1}$$





AdaGrad

Idea of AdaGrad : scale gradients by inverse of square root of sum of historical squares gradients


$$\mathbf{g}_{i+1} = \frac{1}{N_b} \sum_{j=1}^{N_b} \nabla \mathcal{L}_j$$

$$\mathbf{r}_{i+1} = \mathbf{r}_i + \mathbf{g}_{i+1} \cdot \mathbf{g}_{i+1}$$

$$\Delta \boldsymbol{\theta}_{i+1} = -\frac{\alpha}{\delta + \sqrt{\mathbf{r}_{i+1}}} \cdot \mathbf{g}_{i+1}$$

$$\boldsymbol{\theta}_{i+1} = \boldsymbol{\theta}_i + \Delta \boldsymbol{\theta}_{i+1}$$

Parameter dependent rescaling



RMSProp

Idea of RMSProp : gradient accumulation is *exponentially damped*

$$\mathbf{g}_{i+1} = \frac{1}{N_b} \sum_{j=1}^{N_b} \nabla \mathcal{L}_j$$


Slower decay of LR

$$\mathbf{r}_{i+1} = \rho \mathbf{r}_i + (1 - \rho) \mathbf{g}_{i+1} \cdot \mathbf{g}_{i+1}$$

$$\Delta \boldsymbol{\theta}_{i+1} = - \frac{\alpha}{\delta + \sqrt{\mathbf{r}_{i+1}}} \cdot \mathbf{g}_{i+1}$$

$$\boldsymbol{\theta}_{i+1} = \boldsymbol{\theta}_i + \Delta \boldsymbol{\theta}_{i+1}$$

Parameter dependent rescaling



Adam = Momentum + RMSProp

$$\mathbf{g}_{i+1} = \frac{1}{N_b} \sum_{j=1}^{N_b} \nabla \mathcal{L}_j$$

$$\mathbf{m}_{i+1} = \rho_1 \mathbf{m}_i + (1 - \rho_1) \mathbf{g}_{i+1} \leftarrow \text{velocity term}$$

$$\mathbf{v}_{i+1} = \rho_2 \mathbf{v}_i + (1 - \rho_2) \mathbf{g}_{i+1} \cdot \mathbf{g}_{i+1} \leftarrow \text{scaling term}$$


$$\hat{\mathbf{m}}_{i+1} = \frac{\mathbf{m}_{i+1}}{1 - \rho_1^{i+1}} \leftarrow \text{bias correction}$$

$$\hat{\mathbf{v}}_{i+1} = \frac{\mathbf{v}_{i+1}}{1 - \rho_2^{i+1}} \leftarrow \text{bias correction}$$


$$\Delta \boldsymbol{\theta}_{i+1} = - \frac{\alpha}{\delta + \sqrt{\hat{\mathbf{v}}_{i+1}}} \cdot \hat{\mathbf{m}}_{i+1}$$

$$\boldsymbol{\theta}_{i+1} = \boldsymbol{\theta}_i - \Delta \boldsymbol{\theta}_{i+1}$$


$\delta = 10^{-6}$, and two decay rates (ρ_1 and ρ_2).




CNN





Neural Network











Convolutional Neural Networks (CNN)




Convolutional Neural Networks (CNN)



Convolutional Neural Networks (CNN)



Convolution



Mathematical operation on two functions (x, h) that keeps track of the running contributions of the product of these two functions

$$y(t) = \int x(\tau)h(t - \tau)d\tau \leftrightarrow y = x * h$$

Diagram illustrating the convolution operation:

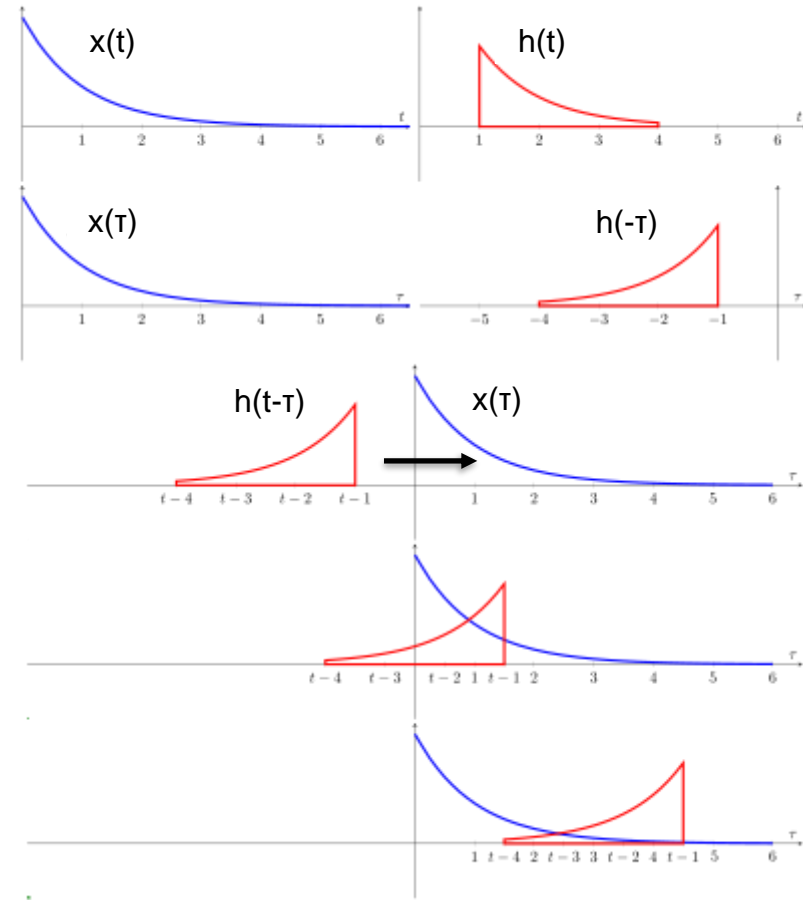
- Output / Feature map (points to $y(t)$)
- Input (points to $x(\tau)$)
- Filter / kernel (points to $h(t - \tau)$)


Signal Proc /
DL

Widely used in signal processing, telecommunication, image processing as it describes the response of a system to an input

Convolution (visual)

$$y(t) = \int x(\tau)h(t - \tau)d\tau \leftrightarrow y = x * h$$

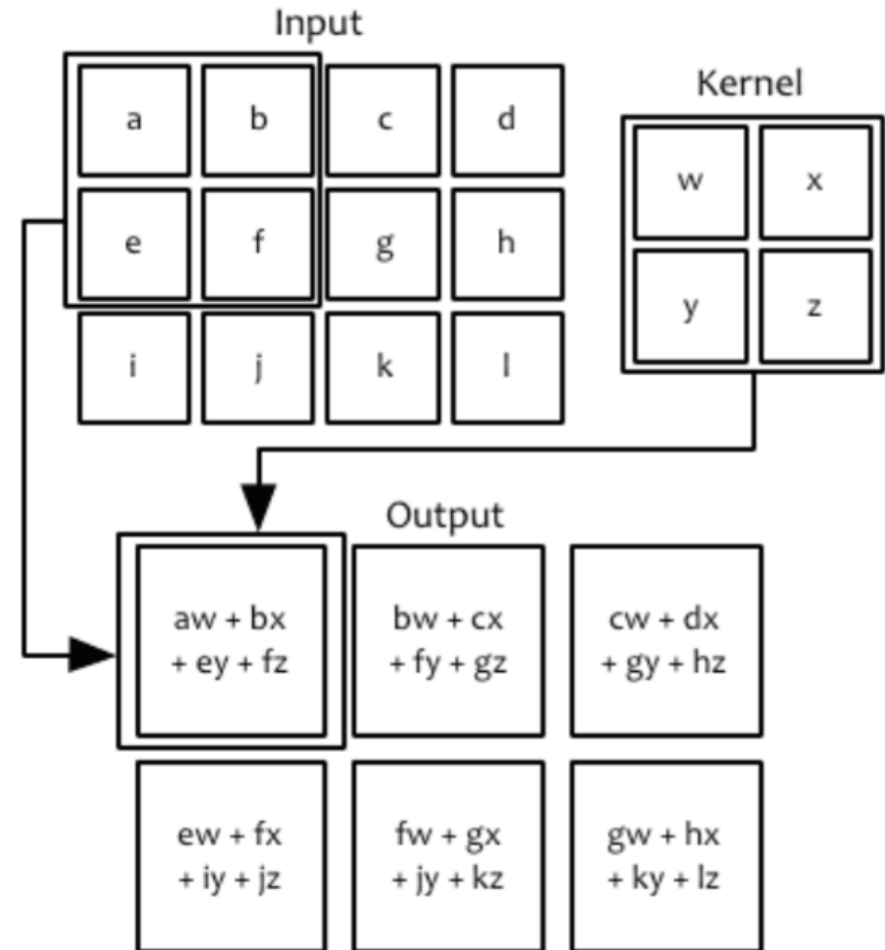




Convolutional Layer


- Input: an image (2-D array): x
- Convolution kernel: W
- Feature map (2-D array of processed data): s
- Convolution operation in 2-D domains:

$$s[i, j] = (x * w)[i, j] = \sum_{m=-M}^M \sum_{n=-N}^N x[i + m, j + n] w[m, n]$$





Convolutional Layer


- Sparse interactions or connectivity / weights;



'90




Horizontal edges



Vertical edges

H filters: hand-crafted

'2000




h_1


h_1

...

h_N





H filters: learned



Convolutional Layer

- Sparse interactions or connectivity / weights;
- Parameter sharing







Convolutional Layer

- Sparse interactions or connectivity / weights;
- Parameter sharing
- Equivariance to translation

shift the input by k samples, the output will also be shifted by the same number of samples;





Convolutional Layer (Kernel)

Filter 1

1	-1	-1
-1	1	-1
-1	-1	1

1	0	0	0	0
0	1	0	0	1
0	0	1	1	0
1	0	0	0	1
0	1	0	0	1
0	0	1	0	1

6 x 6 image

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1


1	-1	-1
-1	1	-1
-1	-1	1

1	0	0	0	0
0	1	0	0	1
0	0	1	1	0
1	0	0	0	1
0	1	0	0	1
0	0	1	0	1

6 x 6 image

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Convolutional Layer (Padding)




2.0 in ₁	3.0 in ₂	0.0 in ₃	5.0 in ₄	2.5 in ₅	0.0 pad ₁
2.0 in ₆	1.5 in ₇	0.5 in ₈	0.0 in ₉	7.0 in ₁₀	0.0 pad ₂
1.5	5.0	5.0	3.0	2.0	0.0
3.0	5.0	7.0	1.5	0.0	0.0
2.0	5.0	2.0	1.5	2.0	0.0
0.0	0.0	0.0	0.0	0.0	0.0



3.0 out ₁	5.0 out ₂	7.0 out ₃

Convolution and Correlation outputs




The size of the output of these two operations can be:


$$N_y = N_x - N_h + 1$$

*when we are interested in the valid signal
(i.e., the entire filter contributes to the output calculation)*


Padding



$$N_y = N_x - N_h + 1$$





No padding



With padding


Strides



With striding


$$N_y = \lfloor (N_x - N_h) / stride + 1 \rfloor.$$

Padding and Strides



$$N_y = \left\lfloor \frac{N_x + 2pad - N_h}{stride} + 1 \right\rfloor$$

Convolutional Layer (Pooling)



- Subsampling pixels will not change the object

bird





Subsampling

bird





We can subsample the pixels to make image smaller,
fewer parameters to characterize the image


Convolutional Layer (Pooling)




Convolutional Layer (Flattening)



Fully Connected
Feedforward network





Convolutional Layer (All in one)




Convolution

Max Pooling

Convolution


Max Pooling

Flattened




cat dog

Channels



In deep learning, we work with multiple feature maps at the same time




Input: can be RGB, spectral bands, horizontal/vertical motion, etc.


Hidden layers: produced by applying multiple filters in previous conv layer

Convolutional layer





As usual, we can apply multiple filters to produce multiple outputs (like in MLP we used to apply multiple weight vectors = weight matrix)









Convolutional Layer (U-NET)





Application

Seismic Data Denoising





Thank You

30 17:50